

Robust Classification via Regression for Learning with Noisy Labels

Sid Bhatia & Aaron Shamouli

Stevens Institute of Technology

December 10, 2024



- ① Introduction
- ② Background: Compositional Data Analysis
- ③ Proposed Method
- ④ Data & Experiments
- ⑤ Code Execution
- ⑥ Other Applications & Analysis
- ⑦ Implementing the Idea

- 1 Introduction
- 2 Background: Compositional Data Analysis
- 3 Proposed Method
- 4 Data & Experiments
- 5 Code Execution
- 6 Other Applications & Analysis
- 7 Implementing the Idea

Background and Motivation

Challenges with Noisy Labels

- Deep neural networks are highly sensitive to noisy labels.
- Noisy labels lead to performance degradation in classification tasks.

Existing Approaches

- **Loss Reweighting:** Focuses on reducing the influence of noisy samples during training.
- **Label Correction:** Attempts to fix noisy labels to improve training stability.

Why This Paper?

Key Motivations

- Existing methods for handling noisy labels are often limited to either loss reweighting or label correction.
- A unified approach combining these strategies could improve classification robustness.

Objective of the Paper

- Develop a method that leverages regression to unify loss reweighting and label correction.
- Apply this method to benchmark datasets with synthetic and real-world noise.

Contributions of the Paper

Key Contributions

- Proposes a novel regression-based framework for classification tasks with noisy labels.
- Demonstrates the effectiveness of the method across synthetic and real-world noisy datasets.
- Bridges the gap between loss reweighting and label correction.

Highlights

- Robust performance on high-noise datasets.
- Outperforms state-of-the-art methods in multiple experiments.

- ① Introduction
- ② Background: Compositional Data Analysis
- ③ Proposed Method
- ④ Data & Experiments
- ⑤ Code Execution
- ⑥ Other Applications & Analysis
- ⑦ Implementing the Idea

Definition and Challenges of Compositional Data

Definition:

- Compositional data are vectors where each element represents a part of a whole (e.g., proportions, percentages).
- Example: $\mathbf{x} = [x_1, x_2, \dots, x_D]$ such that:

$$\sum_{i=1}^D x_i = 1 \quad \text{and} \quad x_i \geq 0, \forall i$$

Challenges:

- Compositional data lie in a constrained simplex, making traditional statistical techniques unsuitable.
- Solutions require transformations that map the data from the simplex to an unconstrained space.

Mapping to an Unconstrained Space

Centered Log-Ratio (clr) Transform:

$$\text{clr}(\mathbf{x}) = \left[\log\left(\frac{x_1}{g(\mathbf{x})}\right), \log\left(\frac{x_2}{g(\mathbf{x})}\right), \dots, \log\left(\frac{x_D}{g(\mathbf{x})}\right) \right]$$

where $g(\mathbf{x}) = (\prod_{i=1}^D x_i)^{\frac{1}{D}}$ is the geometric mean.

Additive Log-Ratio (alr) Transform:

$$\text{alr}(\mathbf{x}) = \left[\log\left(\frac{x_1}{x_D}\right), \log\left(\frac{x_2}{x_D}\right), \dots, \log\left(\frac{x_{D-1}}{x_D}\right) \right]$$

Isometric Log-Ratio (ilr) Transformation

Definition:

- The ilr transformation maps compositional data to an orthonormal basis:

$$\text{ilr}(\mathbf{x}) = \mathbf{V} \cdot \log(\mathbf{x})$$

- Here, \mathbf{V} is a predefined orthonormal basis for the simplex.

Key Property:

- The ilr transform is invertible, enabling mapping back to the original compositional space.
- It is useful for regression and classification tasks.

Advantages of Log-Ratio Transforms

Addressing the Simplex Constraint:

- Transforms map data from the constrained simplex to an unconstrained Euclidean space.
- Enables the application of standard machine learning methods.

Preservation of Ratios:

- Ratios between components are preserved, which is crucial for compositional data analysis.

- ① Introduction
- ② Background: Compositional Data Analysis
- ③ Proposed Method**
- ④ Data & Experiments
- ⑤ Code Execution
- ⑥ Other Applications & Analysis
- ⑦ Implementing the Idea

Three-Step Process for Robust Classification

Key Idea:

- Transform classification into regression by applying log-ratio transformations.
- Incorporate noise modeling and robust regression techniques.
- Map regression outputs back to classification predictions.

Three-Step Process:

- ① Transform classification datasets to regression datasets.
- ② Train using robust regression techniques.
- ③ Convert regression predictions back to classification outputs.

Label Smoothing and Log-Ratio Transform

Label Smoothing:

$$\hat{\mathbf{y}} = (1 - \epsilon) \cdot \mathbf{y} + \epsilon \cdot \frac{\mathbf{1}}{K}$$

where:

- ϵ : Smoothing parameter.
- K : Number of classes.

Log-Ratio Transform:

$$\mathbf{z} = \text{ilr}(\hat{\mathbf{y}})$$

Transforms smoothed classification labels $\hat{\mathbf{y}}$ into an unconstrained Euclidean space.

Handling Noisy Labels with a Gaussian Noise Model

Gaussian Noise Model:

$$\mathbf{z}_{\text{noisy}} = \mathbf{z}_{\text{true}} + \boldsymbol{\eta}, \quad \boldsymbol{\eta} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I})$$

where:

- \mathbf{z}_{true} : True regression targets (from log-ratio transformation).
- $\boldsymbol{\eta}$: Gaussian noise with zero mean and variance σ^2 .

Training Objective:

$$\mathcal{L}_{\text{reg}} = \frac{1}{N} \sum_{i=1}^N \|\mathbf{z}_i - \hat{\mathbf{z}}_i\|_2^2$$

Minimize the mean squared error (MSE) between predicted $\hat{\mathbf{z}}$ and true \mathbf{z} .

Mapping Predictions Back to the Simplex

Inverse Log-Ratio Transform:

$$\hat{\mathbf{y}} = \text{ilr}^{-1}(\hat{\mathbf{z}})$$

Final Prediction:

$$\text{class} = \arg\max_k(\hat{\mathbf{y}}_k)$$

Key Point:

- The predicted $\hat{\mathbf{y}}$ is mapped back to the probability simplex, ensuring valid classification outputs.

Unified Approach for Handling Noisy Labels

Highlights:

- Combines loss reweighting (Gaussian noise modeling) and label correction (log-ratio transformations).
- Ensures robustness by transforming classification into regression.
- Effective on synthetic and real-world noisy datasets.

Advantages:

- Handles high levels of noise effectively.
- Provides interpretable regression-based predictions.

- 1 Introduction
- 2 Background: Compositional Data Analysis
- 3 Proposed Method
- 4 Data & Experiments**
- 5 Code Execution
- 6 Other Applications & Analysis
- 7 Implementing the Idea

Synthetic Datasets

Synthetic Noise:

- **CIFAR-10:**
 - 10 classes, 50,000 training samples, and 10,000 test samples.
 - Noisy labels generated by flipping a percentage of labels.
- **CIFAR-100:**
 - 100 classes, 50,000 training samples, and 10,000 test samples.
 - Higher label complexity with synthetic noise.

Training and Evaluation Details

Model and Training:

- Backbone: WideResNet with depth 28 and width 2.
- Optimizer: Adam with learning rate 0.001.
- Loss Function: Mean squared error (MSE) for regression.

Evaluation Metrics:

- Accuracy: Percentage of correctly classified samples.
- Robustness: Performance under varying noise rates.

Noise Levels Tested:

- Symmetric noise: 20%, 40%, and 60%.
- Asymmetric noise: Realistic noise patterns based on class similarity.

CIFAR-10 and CIFAR-100 Results

CIFAR-10 Results:

- Shifted Gaussian Noise (**SGN**) outperforms baselines at 20%, 40%, and 60% noise levels
- Accuracy improves significantly compared to standard loss reweighting and label correction methods.

CIFAR-100 Results:

- SGN remains robust even with increased class complexity.
- Demonstrates superior performance at high noise levels.

Key Insights

Strengths of SGN:

- Unified approach balances loss reweighting and label correction.
- Consistently outperforms baselines in synthetic noise settings.

Limitations:

- Computational cost is higher due to the regression-based framework.
- Performance may degrade under extreme noise levels ($> 70\%$).

Future Work:

- Explore alternative transformations for compositional data.
- Extend to larger datasets and real-time applications.

- 1 Introduction
- 2 Background: Compositional Data Analysis
- 3 Proposed Method
- 4 Data & Experiments
- 5 Code Execution**
- 6 Other Applications & Analysis
- 7 Implementing the Idea

Our Results vs. SGN

Comparison of Mean Accuracy \pm Standard Deviation:

Method	No Noise (0%)	Symmetric Noise (20%)	Symmetric Noise (40%)	Symmetric Noise (60%)	Asymmetric Noise (20%)	Asymmetric Noise (40%)
CIFAR-10						
SGN	94.12 \pm 0.22	93.02 \pm 0.17	91.29 \pm 0.25	86.03 \pm 1.19	93.35 \pm 0.21	91.26 \pm 0.27
Our Implementation	92.10 \pm 0.25	91.45 \pm 0.20	89.12 \pm 0.30	84.50 \pm 1.10	91.50 \pm 0.23	89.00 \pm 0.25
CIFAR-100						
SGN	73.88 \pm 0.34	71.79 \pm 0.26	66.86 \pm 0.35	56.83 \pm 0.57	72.83 \pm 0.31	71.01 \pm 0.71
Our Implementation	72.10 \pm 0.40	70.00 \pm 0.30	64.80 \pm 0.40	55.00 \pm 0.60	71.00 \pm 0.35	69.50 \pm 0.75

Table 1: Mean Accuracy \pm Standard Deviation for SGN and Our Implementation on CIFAR-10 and CIFAR-100.

Performance on CIFAR-10

Key Observations:

- SGN achieves higher accuracy compared to our implementation across all noise levels.
- Largest accuracy gap is observed under symmetric noise at 40%:

$$\Delta\text{Accuracy} = 91.29 - 89.12 = 2.17\%$$

- Under no noise (0%), our implementation is only 2.02% lower than SGN.
- Both implementations maintain strong performance under asymmetric noise:
 - At 20%, SGN: 93.35%, Ours: 91.50%.
 - At 40%, SGN: 91.26%, Ours: 89.00%.

Strengths of Our Implementation:

- Comparable performance under lower noise levels.
- Slightly lower standard deviations, indicating stable results.

Performance on CIFAR-100

Key Observations:

- SGN performs slightly better than our implementation, especially at higher noise levels:
 - Symmetric noise (60%): SGN: 56.83%, Ours: 55.00%.
 - Asymmetric noise (40%): SGN: 71.01%, Ours: 69.50%.
- Accuracy gap is smaller under no noise:

$$\Delta\text{Accuracy} = 73.88 - 72.10 = 1.78\%$$

Limitations of Our Implementation:

- Larger accuracy gaps at higher noise levels (40% – 60%).
- Higher standard deviations in some cases, indicating less stability.

Comparison Across Both Datasets

General Observations:

- SGN slightly outperforms our implementation across all noise rates and datasets.
- CIFAR-10 results are closer between the two methods than CIFAR-100 results.

Why Does SGN Perform Better?

- Better robustness to high noise levels due to:
 - Advanced loss reweighting strategies.
 - Improved regression-to-classification mapping.
- Possible hyperparameter tuning advantages in SGN.

Future Improvements for Our Implementation:

- Implement better noise modeling techniques (e.g., adaptive noise reweighting).
- Enhance data augmentation to improve generalization under noisy conditions.
- Tune hyperparameters, such as learning rate and model architecture.

- ① Introduction
- ② Background: Compositional Data Analysis
- ③ Proposed Method
- ④ Data & Experiments
- ⑤ Code Execution
- ⑥ Other Applications & Analysis**
- ⑦ Implementing the Idea

Applications of Robust Learning Techniques

- **Medical Diagnosis:** Robust learning can improve classification in medical imaging (e.g., X-rays, MRIs), where mislabeling is common due to human error.
- **Autonomous Vehicles:** Label noise in datasets collected from real-world driving scenarios can impact safety-critical applications.
- **E-Commerce:** Product classification in e-commerce platforms often involves noisy labels, which robust methods can address.
- **Fraud Detection:** Robust regression-based methods can help identify fraudulent transactions by addressing mislabeled data in financial systems.

Analysis of the SGN Approach

- Combines **loss reweighting** and **label correction** effectively, making it robust to varying noise levels.
- Shows significant improvements in datasets like CIFAR-10 and CIFAR-100, even under high symmetric and asymmetric noise rates.
- Compared to baselines like CE and ELR, the SGN approach ensures better generalization, particularly under challenging conditions.

Potential Enhancements to SGN

- **Adaptive Learning Rates:** Experiment with adaptive optimization methods to dynamically adjust learning rates during training.
- **Domain Adaptation:** Extend SGN to handle domain shifts between training and testing datasets.
- **Additional Regularization:** Incorporate dropout or data augmentation techniques to further improve robustness.

- ① Introduction
- ② Background: Compositional Data Analysis
- ③ Proposed Method
- ④ Data & Experiments
- ⑤ Code Execution
- ⑥ Other Applications & Analysis
- ⑦ Implementing the Idea**

Testing the Model on Fashion-MNIST

Why Fashion-MNIST?

- Fashion-MNIST is a drop-in replacement for MNIST with 10 classes of Zalando's article images (e.g., T-shirts, dresses, shoes).
- Contains:
 - 60,000 training samples and 10,000 test samples.
 - Images are grayscale, 28×28 , and labeled with corresponding classes.
- Provides a challenge compared to MNIST due to more complex features and higher inter-class similarity.

Training Details and Evaluation Metrics

Experimental Details:

- Model: WideResNet with depth 28 and width 2.
- Optimizer: Adam with learning rate 0.001.
- Loss Function: Mean squared error (MSE) for regression.
- Noise Levels Tested:
 - Symmetric Noise: 20%, 40%, and 60%.
 - Asymmetric Noise: 20% and 40%.

Evaluation Metrics:

- Accuracy: Percentage of correctly classified samples.
- Standard Deviation: To measure result stability across multiple runs.

Analysis of Results

Key Observations:

- SGN slightly outperforms our implementation, notably under higher noise levels.
- Both methods maintain reasonable performance under no noise, with SGN achieving 91.05% and ours 89.50%.
- Accuracy gaps are more pronounced at 60% symmetric noise:

$$\Delta\text{Accuracy} = 80.20 - 78.00 = 2.20\%$$

Future Directions:

- Investigate architecture adjustments (e.g., deeper WideResNet or additional regularization techniques).
- Explore alternative loss functions to improve robustness under higher noise levels.
- Apply SGN-based models to other datasets, such as SVHN or TinyImageNet.

Thank you for listening!



Sid Bhatia & Aaron Shamouli